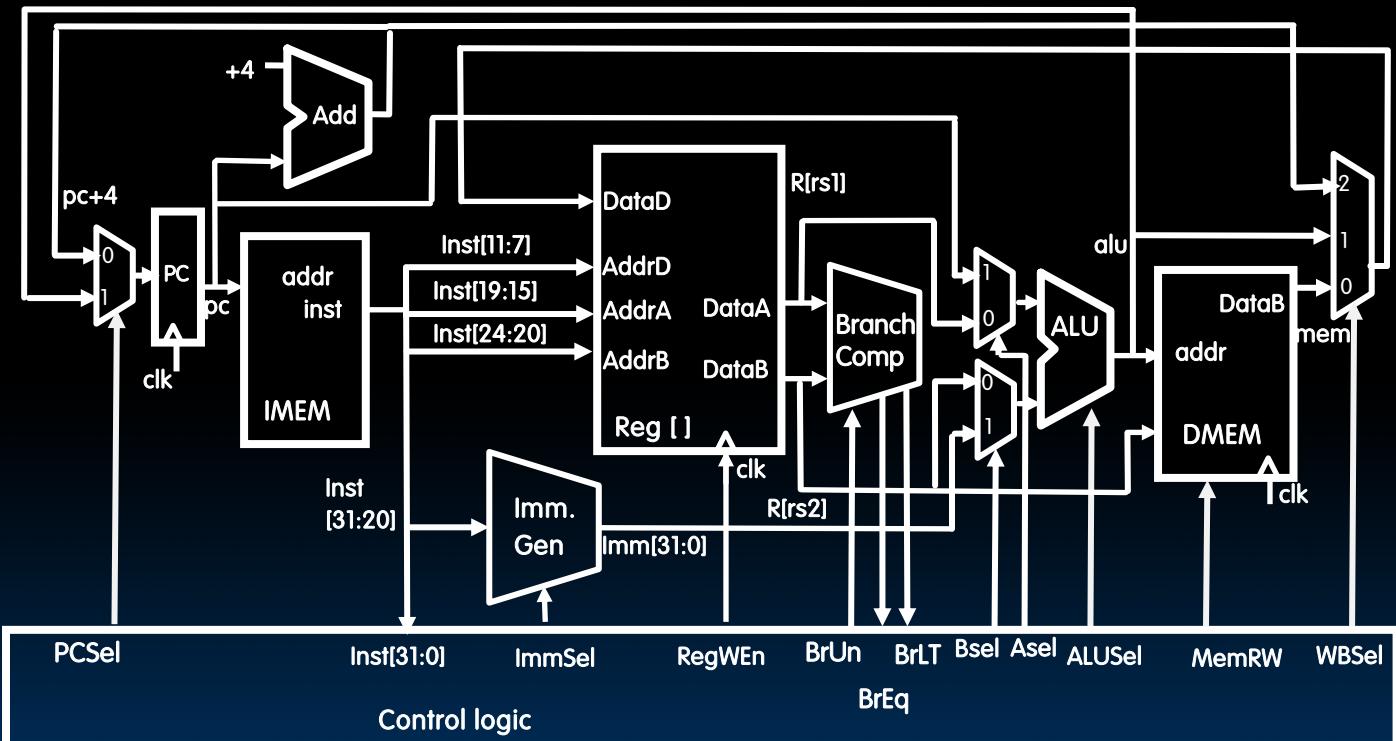


# Control and Status Registers

# Complete Single-Cycle RV32I Datapath!

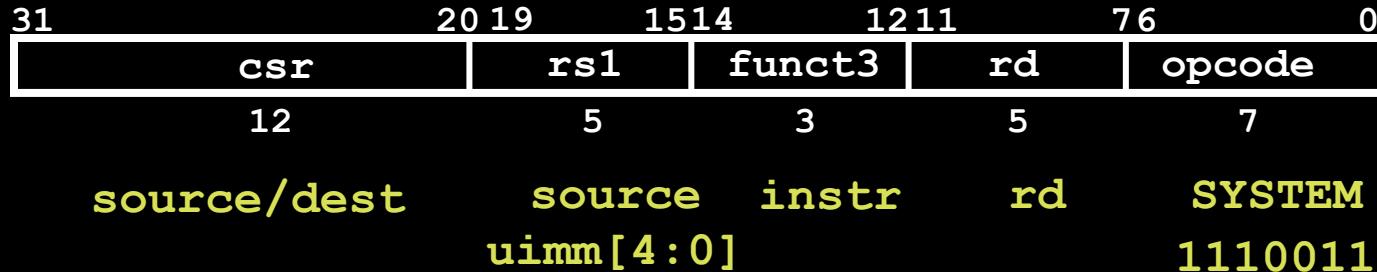




# Control and Status Registers

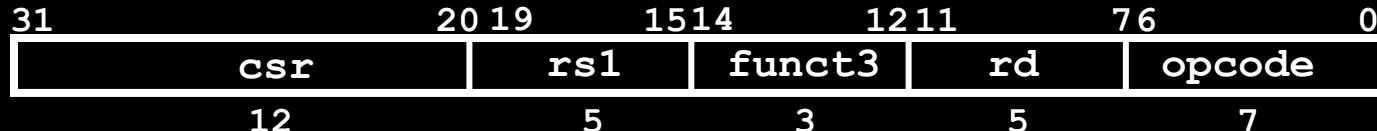
- Control and status registers (CSRs) are separate from the register file (**x0-x31**)
  - Used for monitoring the status and performance
  - There can be up to 4096 CSRs
- Not in the base ISA, but almost mandatory in every implementation
  - ISA is modular
  - Necessary for counters and timers, and communication with peripherals

# CSR Instructions



Register operand				
Instr.	rd	rs	Read CSR?	Write CSR?
<b>csrrw</b>	<b>x0</b>	-	no	yes
<b>csrrw</b>	<b>!x0</b>	-	yes	yes
<b>csrrs/c</b>	-	<b>x0</b>	yes	no
<b>csrrs/c</b>	-	<b>!x0</b>	yes	yes

# CSR Instructions



**source/dest**      **source**    **instr**      **rd**      **SYSTEM**  
uimm[4:0] ← Zero-extended to 32b

Immediate operand				
Instr.	rd	uimm	Read CSR?	Write CSR?
<b>csrrwi</b>	<b>x0</b>	-	no	yes
<b>csrrwi</b>	<b>!x0</b>	-	yes	yes
<b>csrrs/ci</b>	-	0	yes	no
<b>csrrs/ci</b>	-	<b>!0</b>	yes	yes

# CSR Instruction Example

- The CSRRW (Atomic Read/Write CSR) instruction ‘atomically’ swaps values in the CSRs and integer registers.
  - We will see more on ‘atomics’ later
- CSRRW reads the previous value of the CSR and writes it to integer register **rd**. Then writes **rs1** to CSR
- Pseudoinstruction **csrw csr, rs1** is  
**csrrw x0, csr, rs1**
  - **rd=x0**, just writes **rs1** to CSR
- Pseudoinstruction **csrwi csr, uimm** is  
**csrrwi x0, csr, uimm**
  - **rd=x0**, just writes **uimm** to CSR
- Hint: Use write enable and clock...

# System Instructions

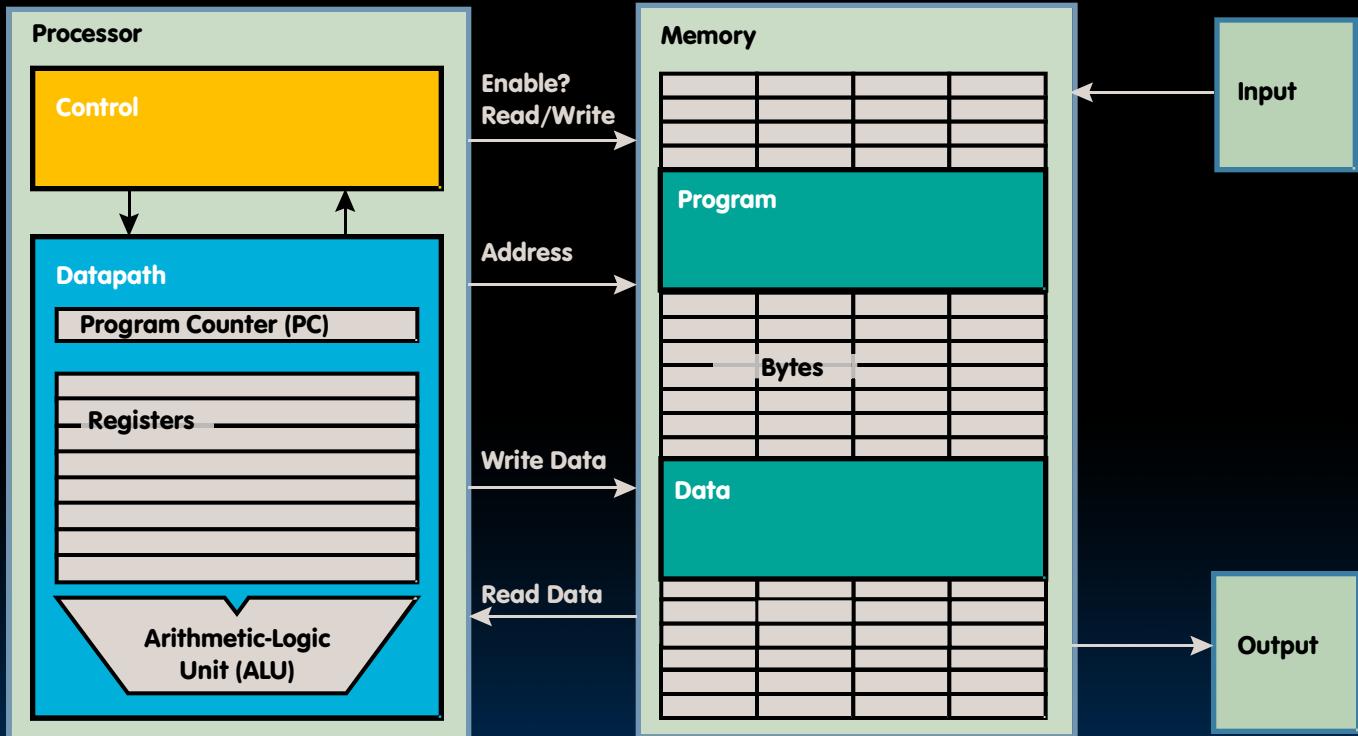
- **ecall** – (I-format) makes requests to supporting execution environment (OS), such as system calls (**syscalls**)
- **ebreak** – (I-format) used e.g. by debuggers to transfer control to a debugging environment

31	20 19	15 14	12 11	7 6	0
funct12	rs1	funct3	rd	opcode	
12	5	3	5	7	
<b>ECALL</b>	0	<b>PRIV</b>	0	<b>SYSTEM</b>	
<b>EBREAK</b>	0	<b>PRIV</b>	0		

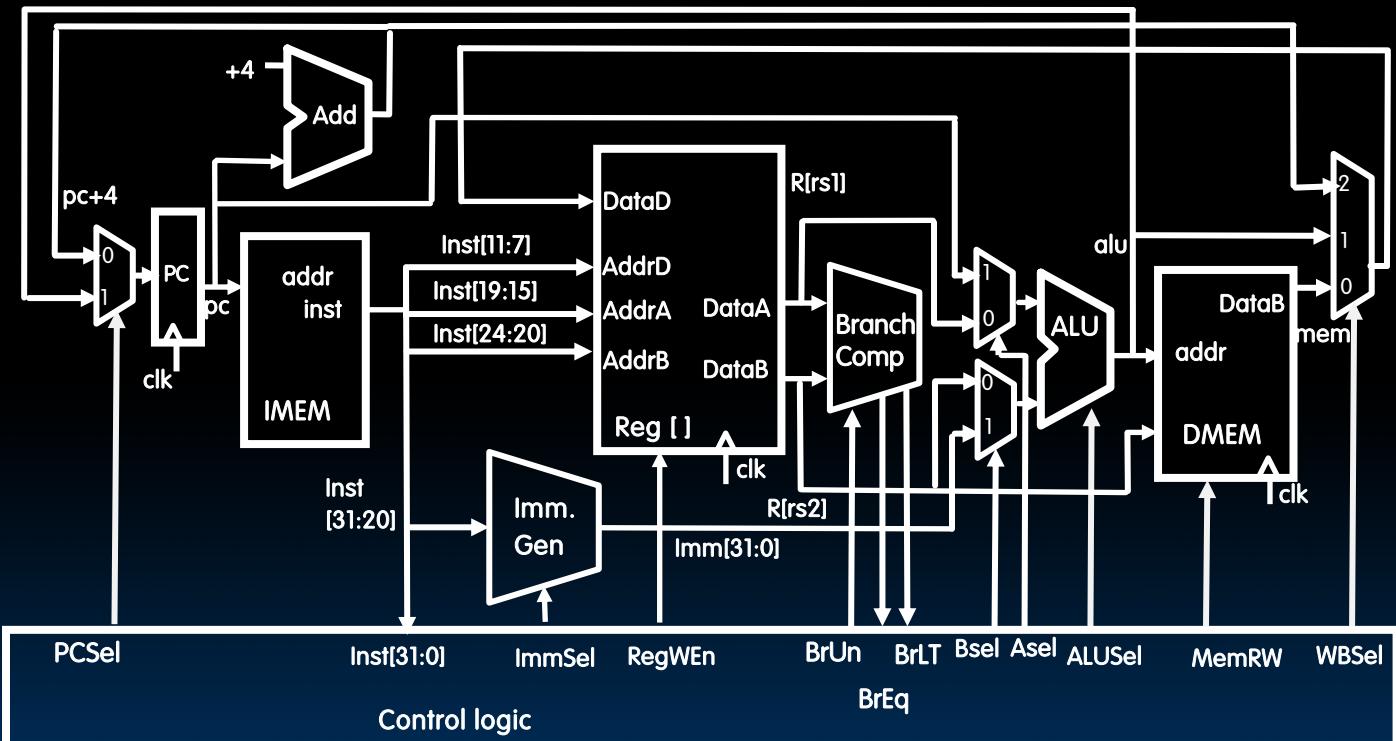
- **fence** – sequences memory (and I/O) accesses as viewed by other threads or co-processors

# Datapath Control

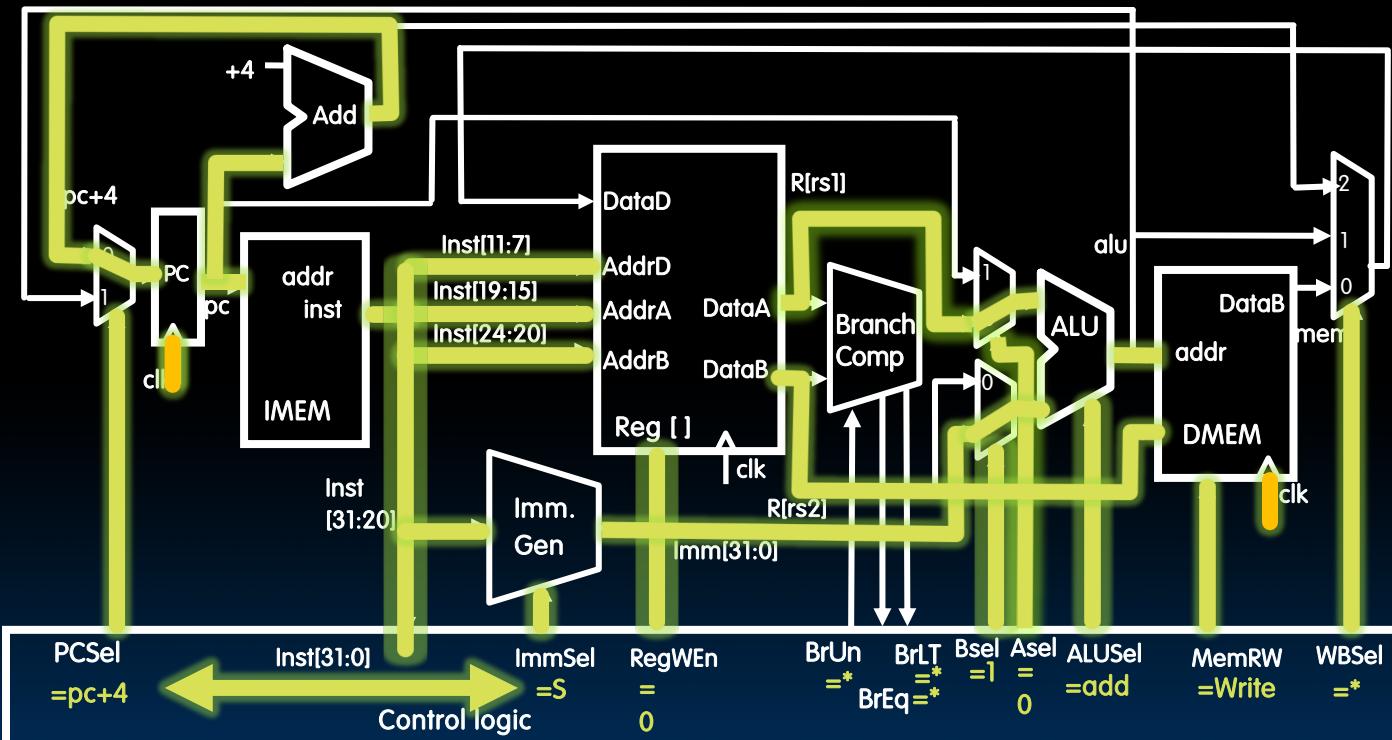
# Our Single-Core Processor



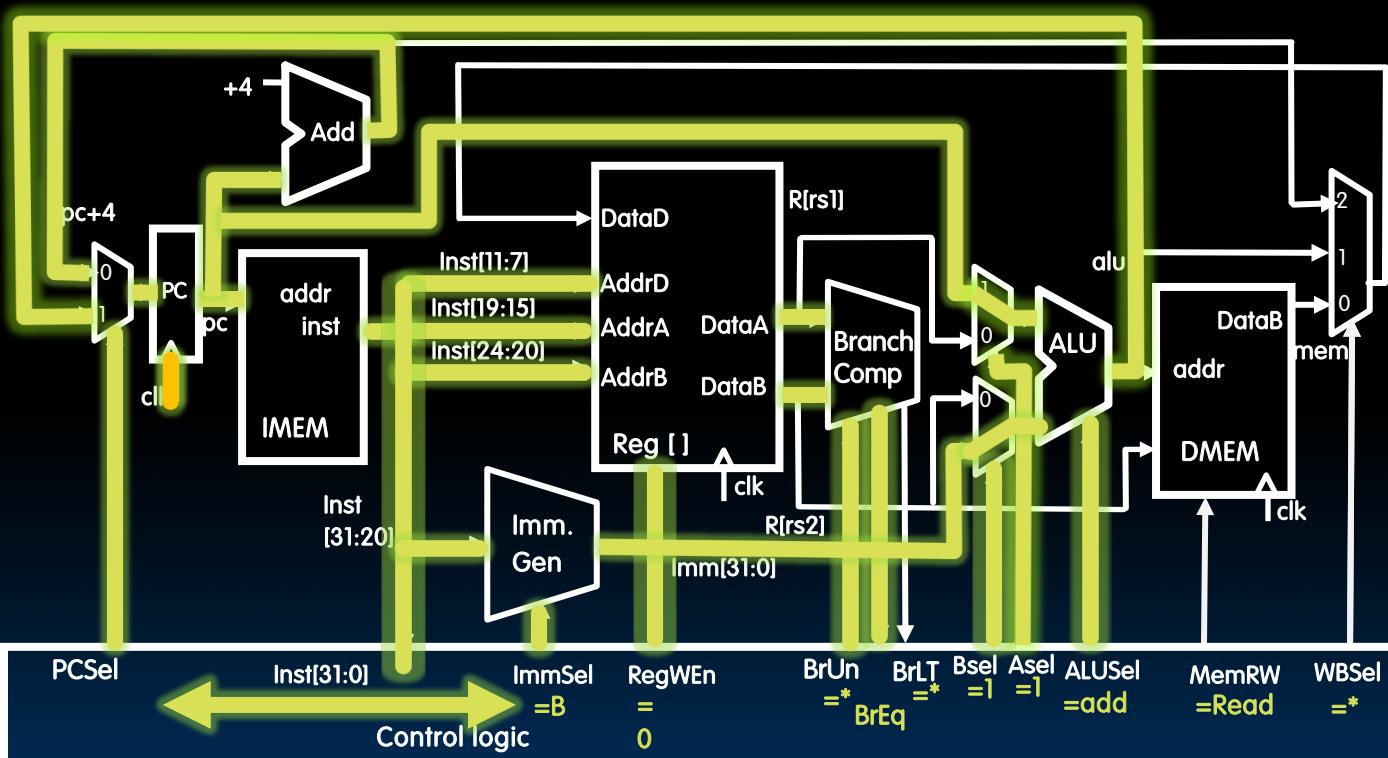
# Single-Cycle RV32I Datapath and Control



# Example: sw

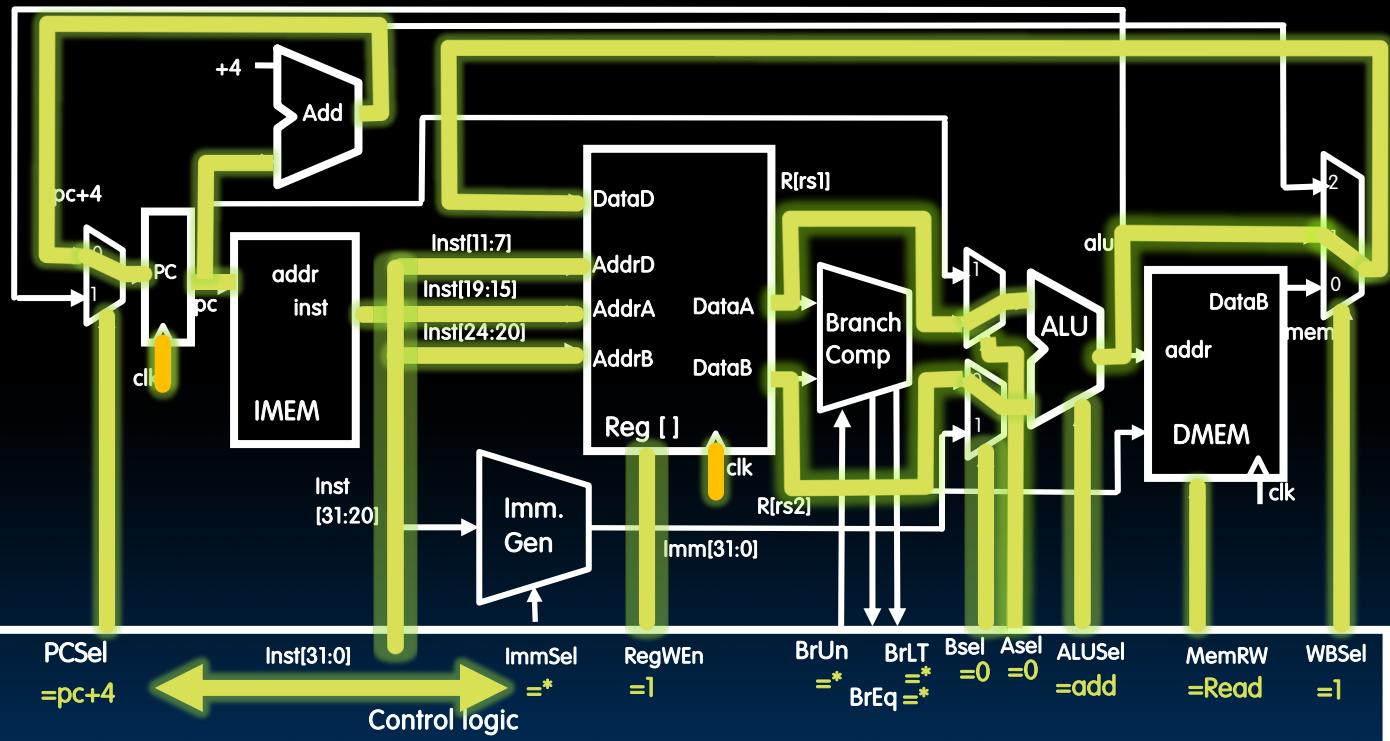


## Example: beq

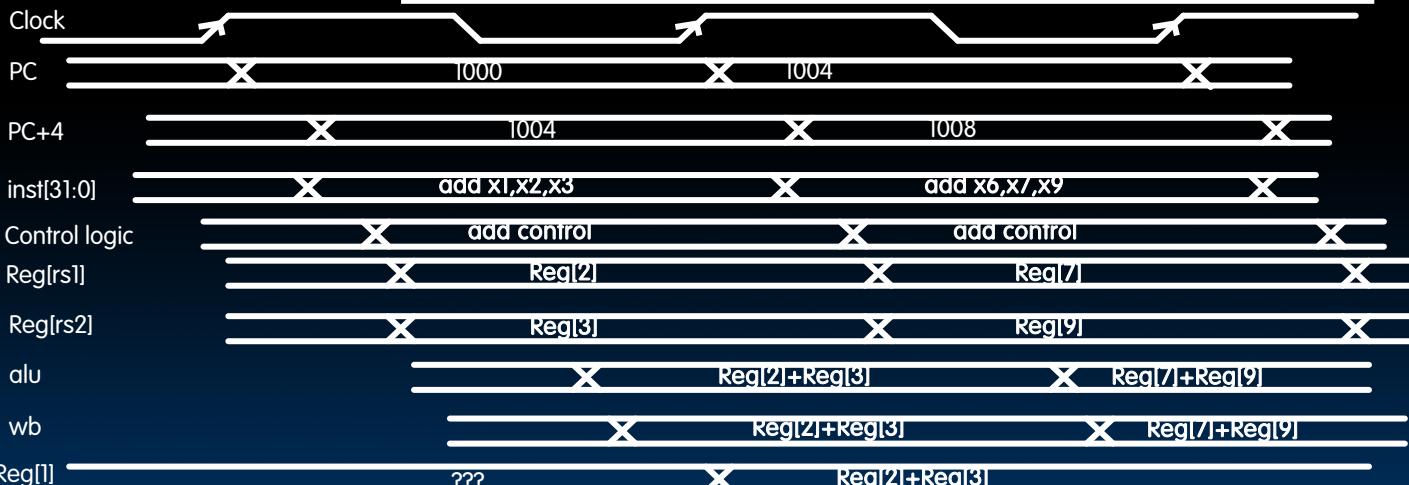
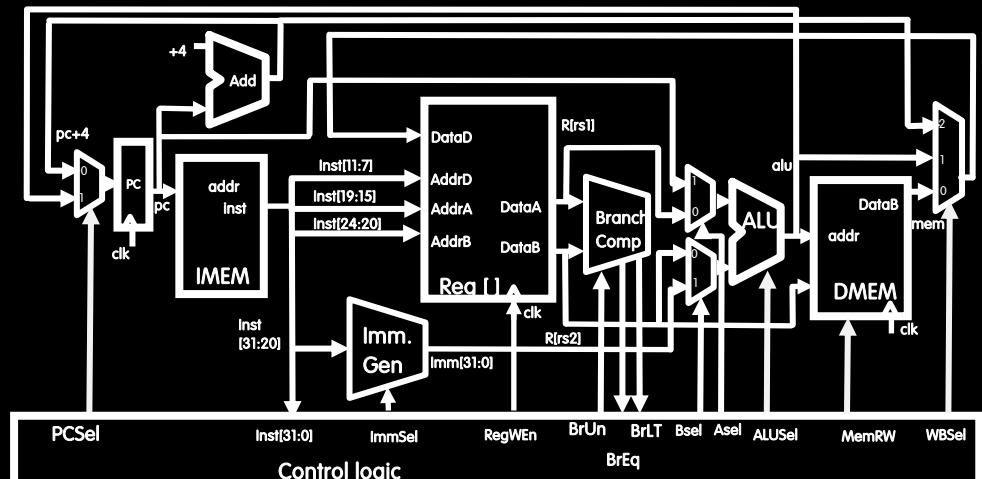


# Instruction Timing

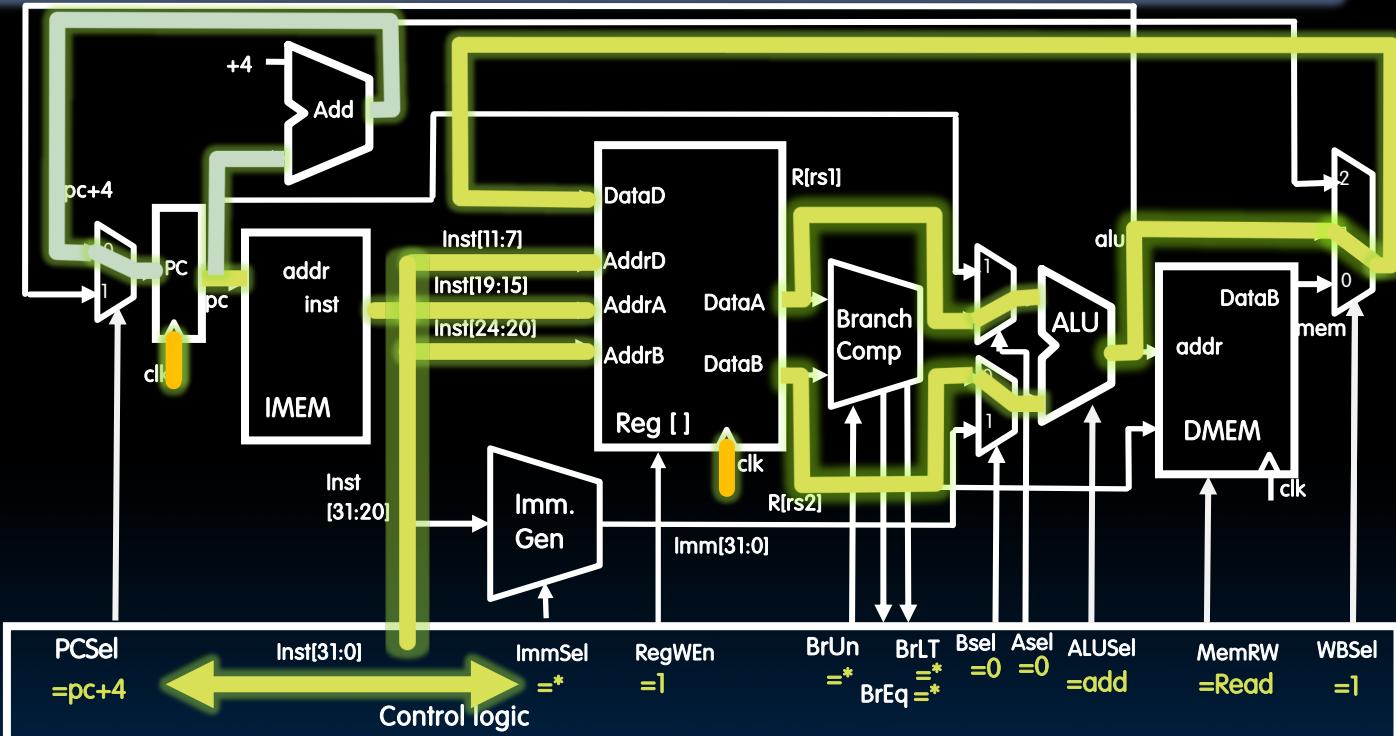
# Example: add



# Add Execution

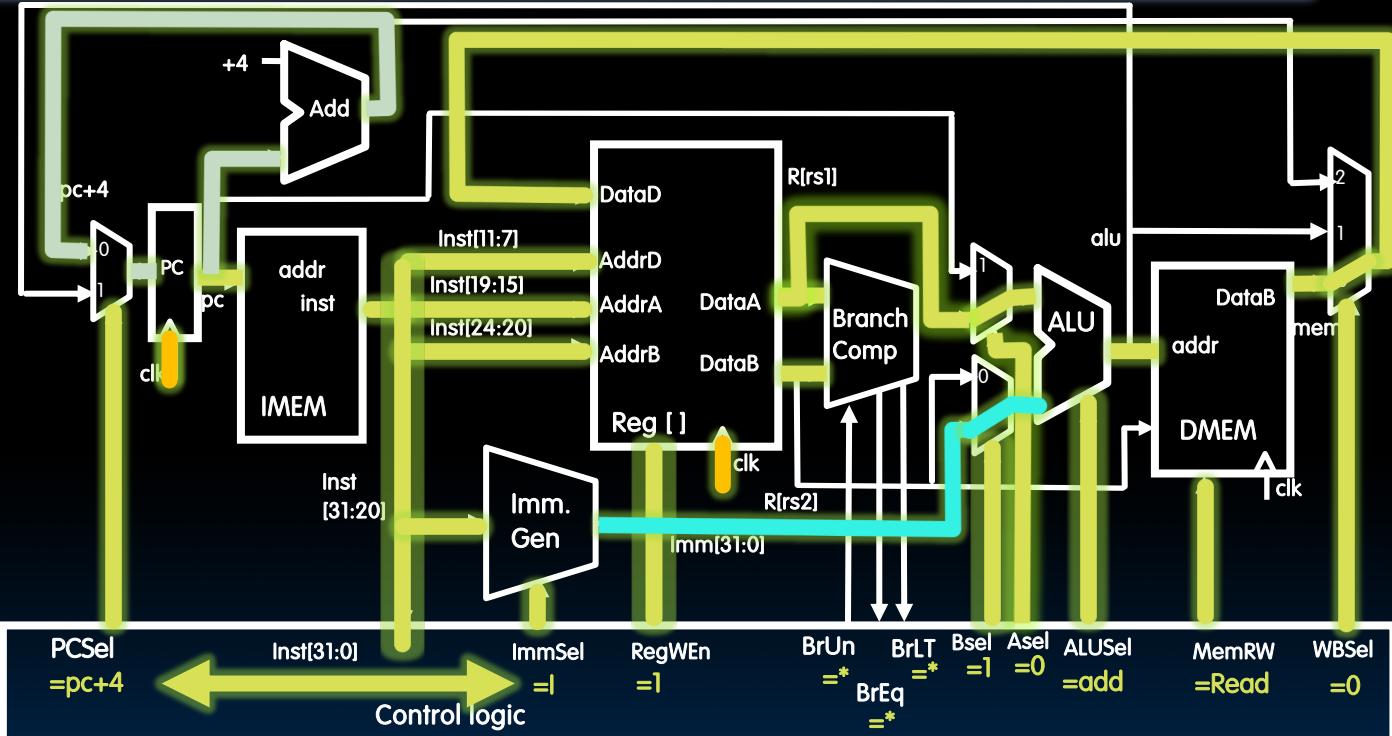


# Example: add timing



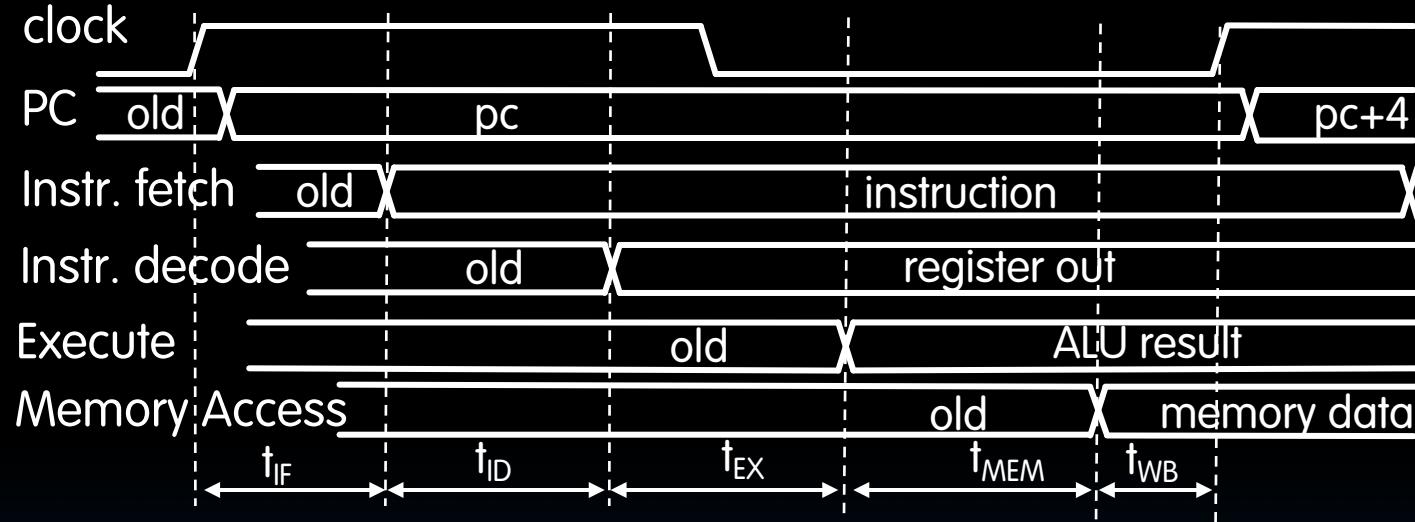
$$\begin{aligned} \text{Critical path} &= t_{\text{clk-q}} + \max \{t_{\text{Add}} + t_{\text{mux}}, t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{mux}} + t_{\text{ALU}} + t_{\text{mux}}\} + t_{\text{setup}} \\ &= t_{\text{clk-q}} + t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{mux}} + t_{\text{ALU}} + t_{\text{mux}} + t_{\text{setup}} \end{aligned}$$

## Example: lw



Critical path =  $t_{clk-q} + \max \{t_{Add} + t_{mux}, t_{IMEM} + t_{Imm} + t_{mux} + t_{ALU} + t_{DMEM} + t_{mux}, t_{IMEM} + t_{Reg} + t_{mux} + t_{ALU} + t_{DMEM} + t_{mux}\} + t_{setup}$

# Instruction Timing



IF	ID	EX	MEM	WB	Total
I-MEM	Reg Read	ALU	D-MEM	Reg W	
200 ps	100 ps	200 ps	200 ps	100 ps	800 ps

# Instruction Timing

Instr	IF = 200ps	ID = 100ps	ALU = 200ps	MEM=200ps	WB = 100ps	Total
add	X	X	X		X	600ps
beq	X	X	X			500ps
jal	X	X	X		X	600ps
lw	X	X	X	X	X	800ps
sw	X	X	X	X		700ps

- Maximum clock frequency
  - $f_{max} = 1/800\text{ps} = 1.25 \text{ GHz}$
- Most blocks idle most of the time
  - E.g.  $f_{max,ALU} = 1/200\text{ps} = 5 \text{ GHz!}$

# Control Logic Design



# Control Logic Truth Table

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
<b>add</b>	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
<b>sub</b>	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
<i>(R-R Op)</i>	*	*	+4	*	*	Reg	Reg	<i>(Op)</i>	Read	1	ALU
<b>addi</b>	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
<b>lw</b>	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
<b>sw</b>	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
<b>beq</b>	0	*	+4	B	*	PC	Imm	Add	Read	0	*
<b>beq</b>	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
<b>bne</b>	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
<b>bne</b>	1	*	+4	B	*	PC	Imm	Add	Read	0	*
<b>blt</b>	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
<b>bltu</b>	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
<b>jalr</b>	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
<b>jal</b>	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
<b>auipc</b>	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU

# Control Realization Options

- ROM
  - "Read-Only Memory"
  - Regular structure
  - Can be easily reprogrammed
    - fix errors
    - add instructions
  - Popular when designing control logic manually
- Combinatorial Logic
  - Today, chip designers use logic synthesis tools to convert truth tables to networks of gates

# RV32I, A Nine-Bit ISA!

imm[31:12]			rd	011011	LUI
imm[31:12]			rd	001011	AUIPC
imm[20:10:11 11:19:12]			rd	110111	JAL
imm[11:0]		rs1	000	imm[4:1 11]	JALR
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	BEQ
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	BNE
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	BLT
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	BGE
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	BLTU
imm[11:0]		rs1	000	rd	BGEU
imm[11:0]		rs1	001	rd	LB
imm[11:0]		rs1	010	rd	LH
imm[11:0]		rs1	100	rd	LW
imm[11:0]		rs1	101	rd	LBU
imm[11:0]		rs1	000	imm[4:0]	LHU
imm[11:5]	rs2	rs1	001	imm[4:0]	SB
imm[11:5]	rs2	rs1	010	imm[4:0]	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	SW
imm[11:0]		rs1	000	rd	ADDI
imm[11:0]		rs1	010	rd	SLTI
imm[11:0]		rs1	011	rd	SLTIU
imm[11:0]		rs1	100	rd	XORI
imm[11:0]		rs1	110	rd	ORI
imm[11:0]		rs1	111	rd	ANDI
000000	shamt	rs1	001	rd	SLLI
000000	shamt	rs1	101	rd	SRLI
000000	shamt	rs1	101	rd	SRAI
000000	rs2	rs1	000	rd	ADD
000000	rs2	rs1	000	rd	SUB
000000	rs2	rs1	001	rd	SU
000000	rs2	rs1	010	rd	SLT
000000	rs2	rs1	011	rd	SLTU
000000	rs2	rs1	100	rd	XOR
000000	rs2	rs1	101	rd	SRL
000000	rs2	rs1	101	rd	SRA
000000	rs2	rs1	110	rd	OR
000000	rs2	rs1	111	rd	AND
fm	pred	succ	rs1	000	FENCE
000000000000			00000	00000	ECALL
000000000001			00000	00000	EBREAK

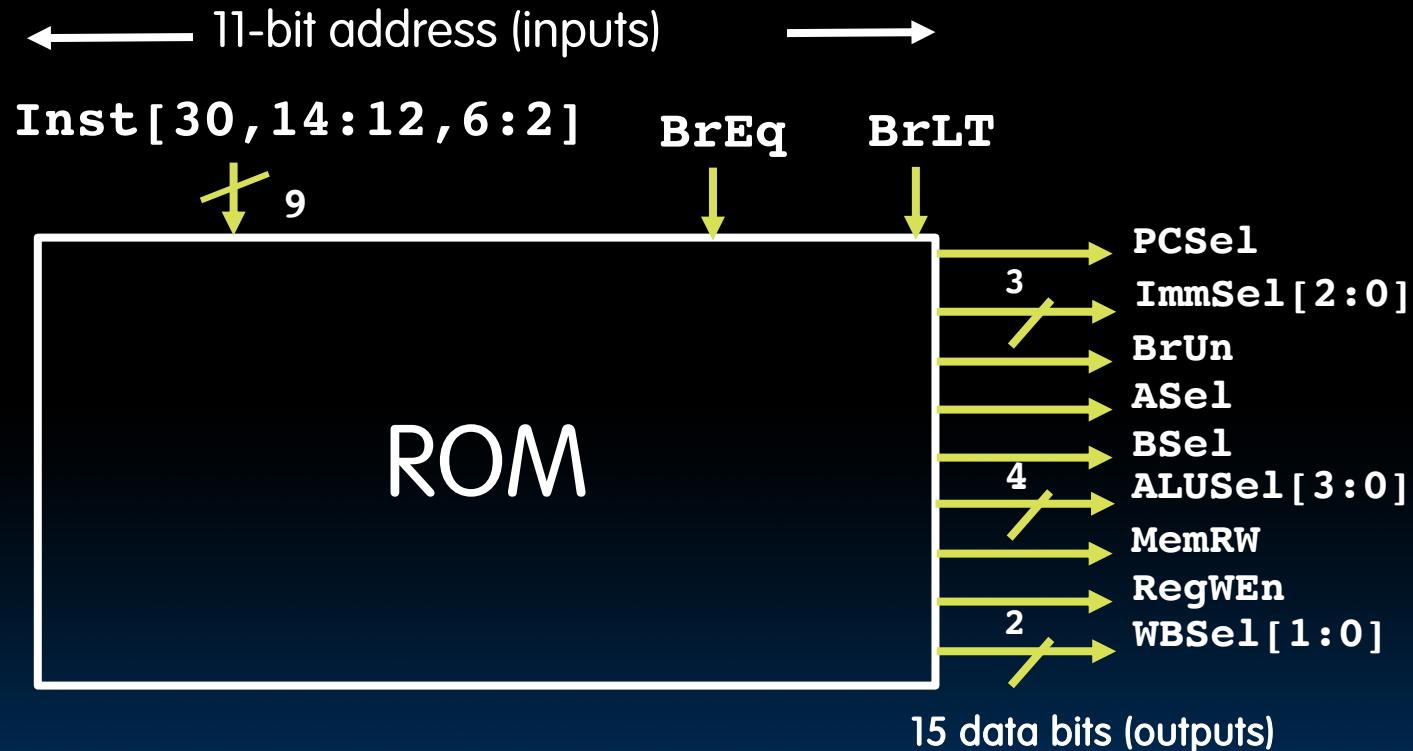
- Instruction type encoded using only 9 bits:
- inst[30], inst[14:12], inst[6:2]**

**inst[6:2]**

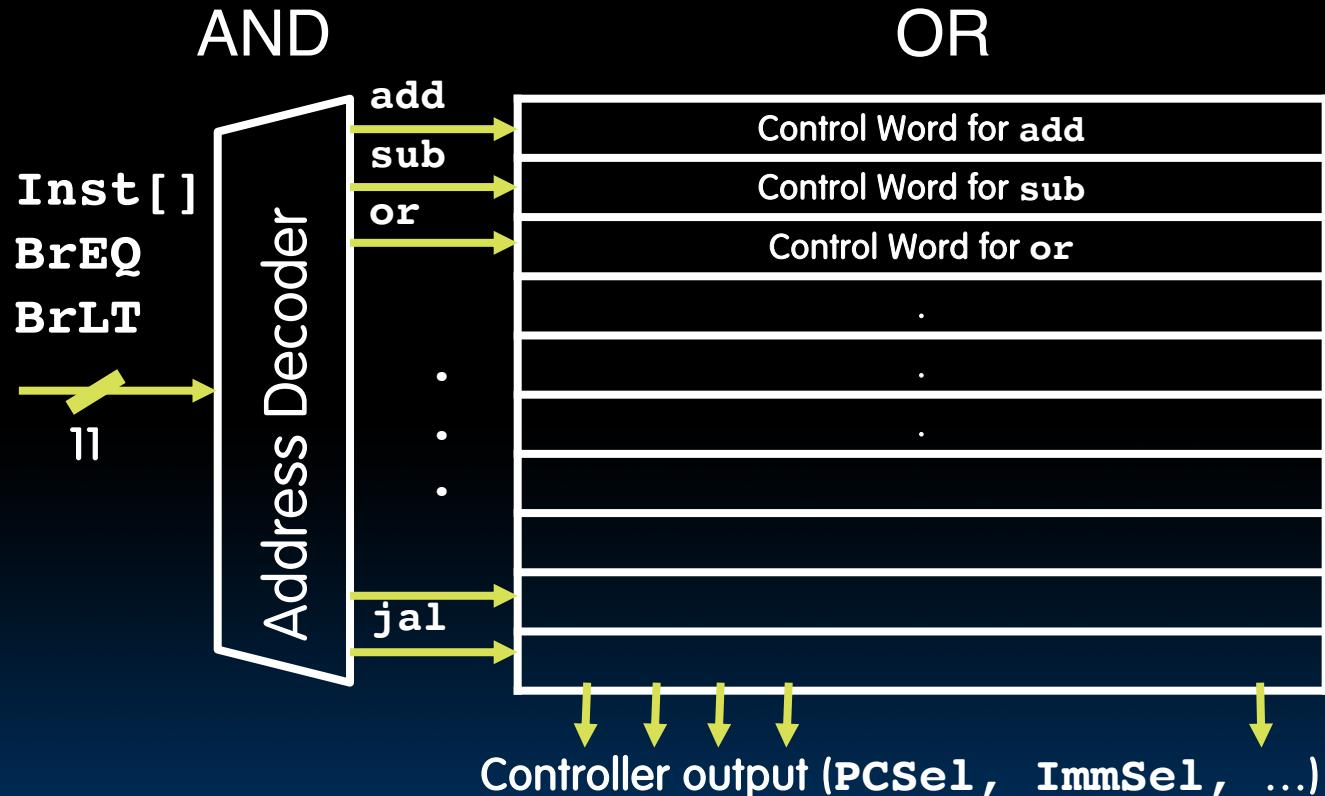
**inst[14:12]**

**inst[30]**

# ROM-based Control



# ROM Controller Implementation



# Combinational Logic Control

- Simplest example: **BrUn**

	inst[14:12]	inst[6:2]	
imm[11:2][10:5]	rs2	rs1	000      imm[4:1][11]      1100011      BEQ
imm[11:2][10:5]	rs2	rs1	001      imm[4:1][11]      1100011      BNE
imm[11:2][10:5]	rs2	rs1	100      imm[4:1][11]      1100011      BLT
imm[11:2][10:5]	rs2	rs1	101      imm[4:1][11]      1100011      BGE
imm[11:2][10:5]	rs2	rs1	110      imm[4:1][11]      1100011      BLTU
imm[11:2][10:5]	rs2	rs1	111      imm[4:1][11]      1100011      BGEU

- How to decode whether BrUn is 1?

$$\text{BrUn} = \text{Inst}[13] \bullet \text{Branch}$$

# Control Logic to Decode add

inst[30]	inst[14:12]	inst[6:2]	
000000	shamt	rs1	001 SLLI
000000	shamt	rs1	101 SRLI
100000	shamt	rs1	101 SRAI
000000	rs2	rs1	000 ADD
100000	rs2	rs1	000 SUB
000000	rs2	rs1	001 SLL
000000	rs2	rs1	010 SLT
000000	rs2	rs1	011 SLTU
000000	rs2	rs1	100 XOR
000000	rs2	rs1	101 SRL
100000	rs2	rs1	101 SRA
000000	rs2	rs1	110 OR
000000	rs2	rs1	111 AND

$$\text{add} = \overline{i[30]} \cdot \overline{i[14]} \cdot \overline{i[13]} \cdot \overline{i[12]} \cdot \text{R-type}$$

$$\text{R-type} = \overline{i[6]} \cdot \overline{i[5]} \cdot \overline{i[4]} \cdot \overline{i[3]} \cdot \overline{i[2]} \cdot \text{RV32I}$$

$$\text{RV32I} = i[1] \cdot i[0]$$

**“And In  
Conclusion...”**

# Call home, we've made HW/SW contact!

High Level Language  
Program (e.g., C)

| Compiler

Assembly Language  
Program (e.g., RISC-V)

| Assembler

Machine Language  
Program (RISC-V)

```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;
```

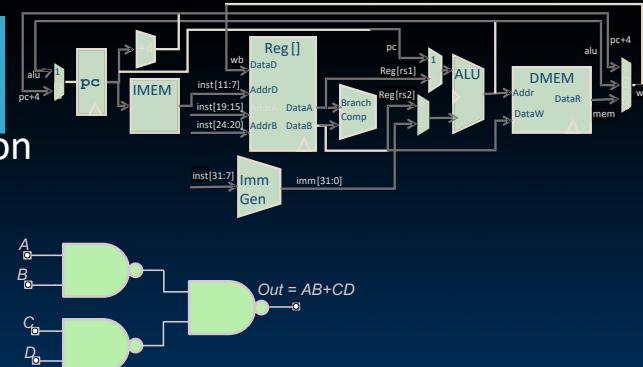
```
lw    x3, 0(x10)  
lw    x4, 4(x10)  
sw    x4, 0(x10)  
sw    x3, 4(x10)
```

1000	1101	1110	0010	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000
1000	1110	0001	0000	0000	0000	0000	0000	0000	0000	0000	0000	0100			
1010	1110	0001	0010	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000	0000
1010	1101	1110	0010	0000	0000	0000	0000	0000	0000	0000	0000	0100			

Hardware Architecture Description  
(e.g., block diagrams)

| Architecture Implementation

Logic Circuit Description  
(Circuit Schematic Diagrams)





# “And In conclusion...”

---

- We have built a processor!
  - Capable of executing all RISC-V instructions in one cycle each
  - Not all units (hardware) used by all instructions
  - Critical path changes
- 5 Phases of execution
  - IF, ID, EX, MEM, WB
  - Not all instructions are active in all phases
- Controller specifies how to execute instructions
  - Implemented as ROM or logic



UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in **Computer Architecture** (a.k.a. Machine Structures)



UC Berkeley  
Professor  
Bora Nikolić

## Pipelining

# New-School Machine Structures

## Software

Parallel Requests

Assigned to computer

e.g., Search "Cats"

Parallel Threads

Assigned to core e.g., Lookup, Ads

**Parallel Instructions**

>1 instruction @ one time

e.g., 5 pipelined instructions

Parallel Data

>1 data item @ one time

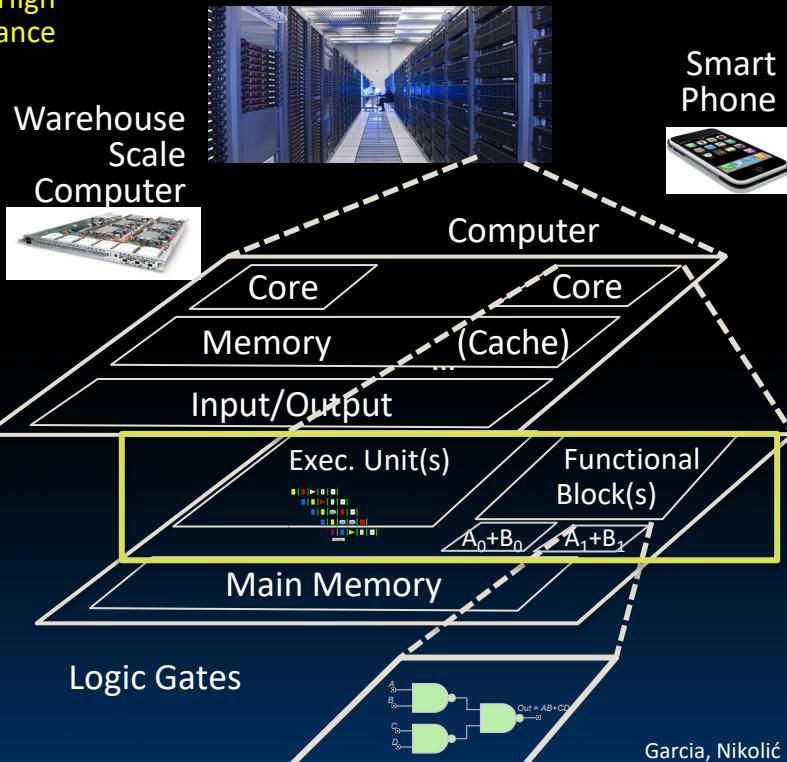
e.g., Add of 4 pairs of words

Hardware descriptions

All gates work in parallel at same time

Harness  
Parallelism &  
Achieve High  
Performance

## Hardware



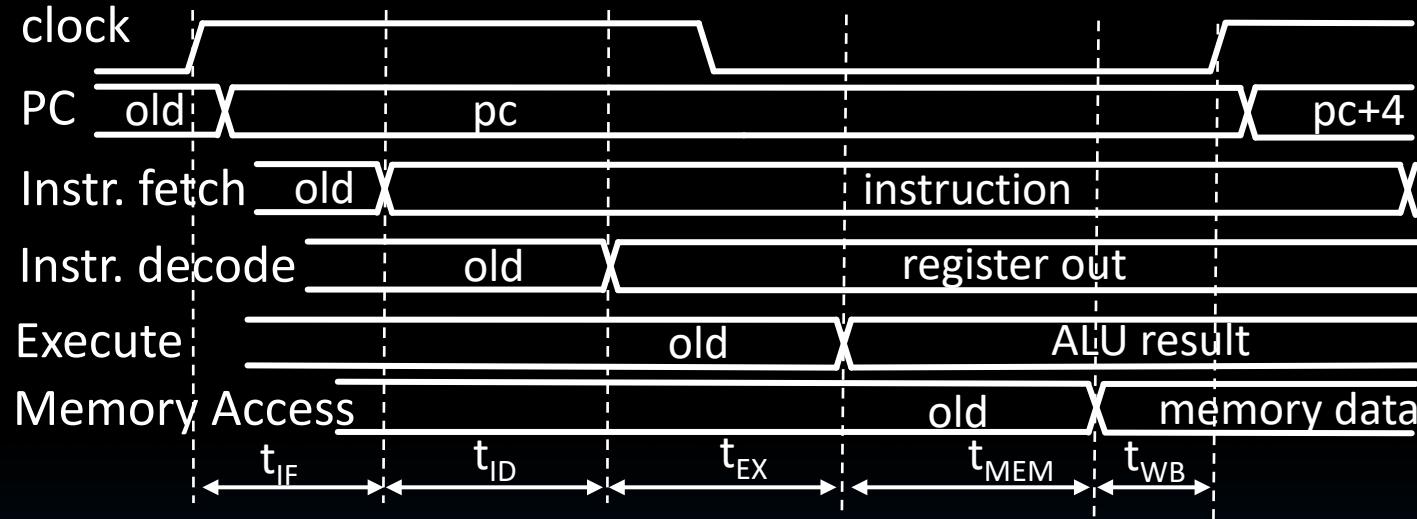


# 6 Great Ideas in Computer Architecture

---

1. Abstraction (Layers of Representation/Interpretation)
2. Moore's Law
3. Principle of Locality/Memory Hierarchy
4. Parallelism
5. Performance Measurement & Improvement
6. Dependability via Redundancy

# Instruction Timing



IF	ID	EX	MEM	WB	Total
I-MEM	Reg Read	ALU	D-MEM	Reg W	
200 ps	100 ps	200 ps	200 ps	100 ps	<b>800 ps</b>

# Instruction Timing

Instr	IF = 200ps	ID = 100ps	ALU = 200ps	MEM=200ps	WB = 100ps	Total
add	X	X	X		X	600ps
beq	X	X	X			500ps
jal	X	X	X		X	600ps
lw	X	X	X	X	X	800ps
sw	X	X	X	X		700ps

- Maximum clock frequency
  - $f_{max} = 1/800\text{ps} = 1.25 \text{ GHz}$

# Performance Measures

- “Our” Single-cycle RISC-V CPU executes instructions at 1.25 GHz
  - 1 instruction every 800 ps
- Can we improve its performance?
  - What do we mean with this statement?
  - Not so obvious:
    - Quicker response time, so one job finishes faster?
    - More jobs per unit time (e.g. web server returning pages, spoken words recognized)?
    - Longer battery life?

# Transportation Analogy



	Sports Car	Bus
Passenger Capacity	2	50
Travel Speed	200 mph	50 mph
Gas Mileage	5 mpg	2 mpg

**50 Mile trip (assume they return instantaneously)**

	Sports Car	Bus
Travel Time	15 min	60 min
Time for 100 passengers	750 min (50 2-person trips)	120 min (two 50-person trips)
Gallons per passenger	5 gallons	0.5 gallons

# Computer Analogy

Transportation	Computer
Trip Time	Program execution time: e.g. time to update display
Time for 100 passengers	Throughput: e.g. number of server requests handled per hour
Gallons per passenger	Energy per task*: e.g. how many movies you can watch per battery charge or energy bill for datacenter

\* Note: Power is not a good measure, since low-power CPU might run for a long time to complete one task consuming more energy than faster computer running at higher power for a shorter time

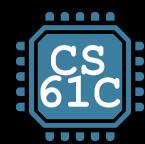
# Processor Performance Iron Law

# “Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$



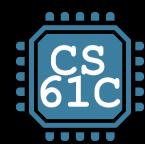
**CPI = Cycles Per Instruction**



# Instructions per Program

Determined by

- Task
- Algorithm, e.g.  $O(N^2)$  vs  $O(N)$
- Programming language
- Compiler
- Instruction Set Architecture (ISA)



# (Average) Clock Cycles per Instruction (CPI)

## Determined by

- ISA
- Processor implementation (or *microarchitecture*)
- E.g. for “our” single-cycle RISC-V design, CPI = 1
- Complex instructions (e.g. **strcpy**), CPI  $\gg$  1
- Superscalar processors, CPI  $<$  1 (next lectures)

# Time per Cycle (1/Frequency)

## Determined by

- Processor microarchitecture (determines critical path through logic gates)
- Technology (e.g. 5nm versus 28nm)
- Power budget (lower voltages reduce transistor speed)

# Speed Tradeoff Example

- For some task (e.g. image compression) ...

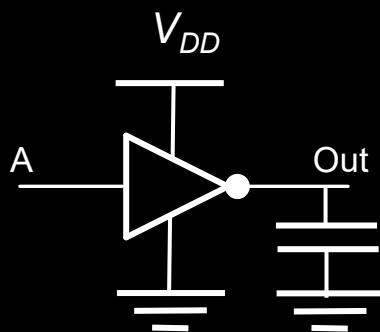
	Processor A	Processor B
# Instructions	1 Million	1.5 Million
Average CPI	2.5	1
Clock rate $f$	2.5 GHz	2 GHz
Execution time	1 ms	0.75 ms

Processor B is faster for this task, despite executing more instructions and having a slower clock rate!

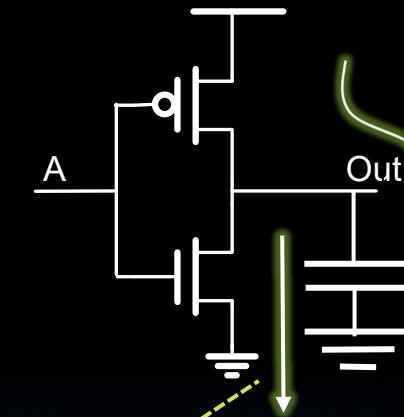
**Energy  
Efficiency**

# Where Does Energy Go in CMOS?

Symbol (INV)



Schematic



Leakage

Charging  
capacitors  
( $CV^2$ )

(30%)

(70%)

# Energy per Task

$$\frac{\text{Energy}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Energy}}{\text{Instruction}}$$

$$\frac{\text{Energy}}{\text{Program}} \propto \frac{\text{Instructions}}{\text{Program}} * C V^2$$

“Capacitance” depends on technology, processor features e.g. # of cores

Supply voltage, e.g. 1V

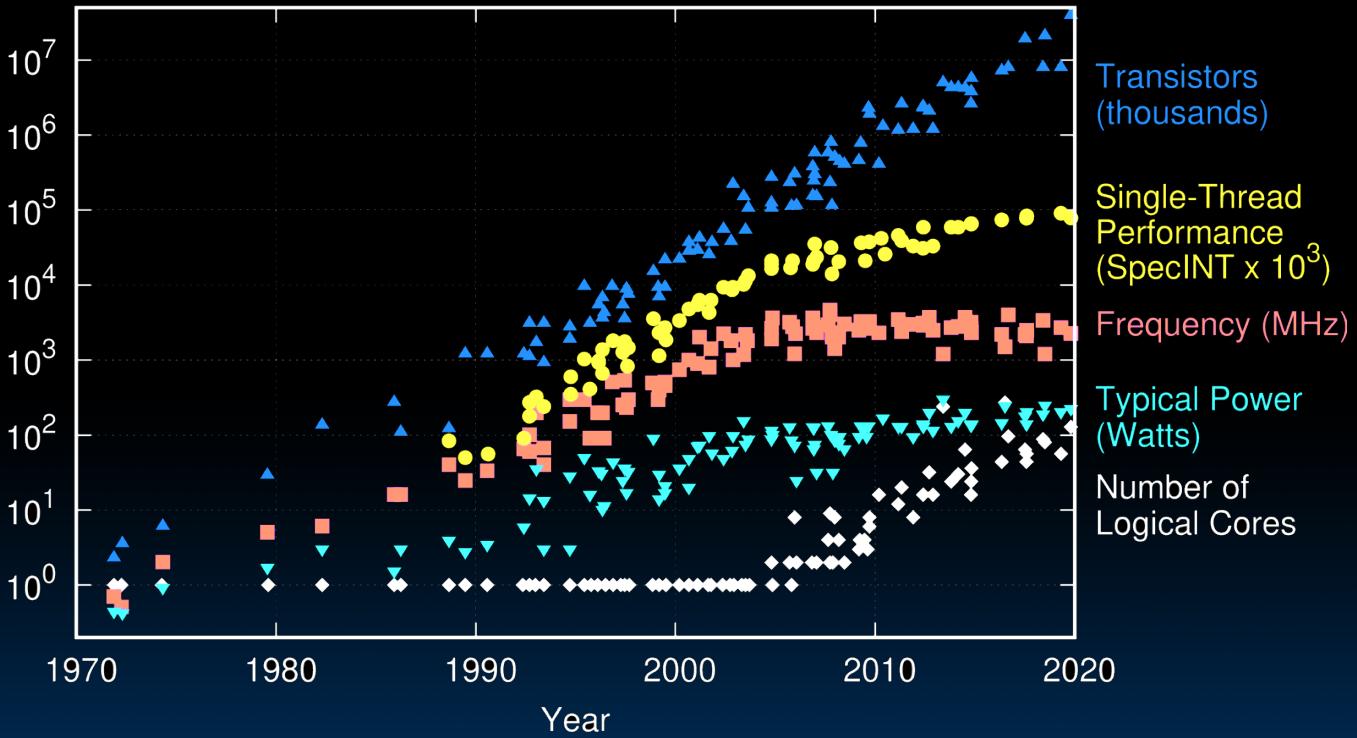
Want to reduce capacitance and voltage to reduce energy/task

# Energy Tradeoff Example

- “Next-generation” processor
  - C (Moore’s Law): -15 %
  - Supply voltage,  $V_{\text{sup}}$ : -15 %
  - Energy consumption:  $0 - (1 - 0.85^3) = \text{-39 \%}$
- Significantly improved energy efficiency thanks to
  - Moore’s Law AND
  - Reduced supply voltage

# Performance/Power Trends

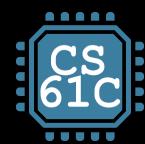
48 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2019 by K. Rupp

# End of Scaling

- In recent years, industry has not been able to reduce supply voltage much, as reducing it further would mean increasing “leakage power” where transistor switches don’t fully turn off (more like dimmer switch than on-off switch)
- Also, size of transistors and hence capacitance, not shrinking as much as before between transistor generations
  - Need to go to 3D
- Power becomes a growing concern – the “power wall”



# Energy “Iron Law”

$$\text{Performance} = \frac{\text{Power}}{(\text{Tasks}/\text{Second})} * \frac{\text{Energy Efficiency}}{(\text{Joules}/\text{Second})}$$

- Energy efficiency (e.g., instructions/Joule) is key metric in all computing devices
- For power-constrained systems (e.g., 20MW datacenter), need better energy efficiency to get more performance at same power
- For energy-constrained systems (e.g., 1W phone), need better energy efficiency to prolong battery life

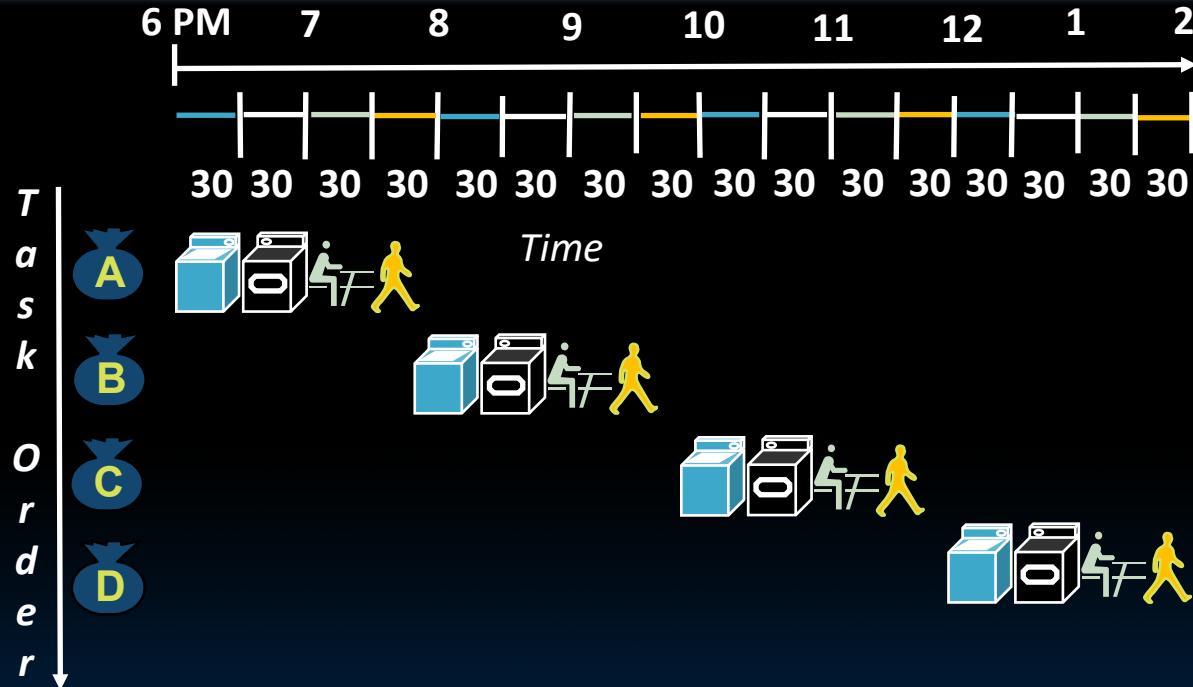
# Introduction to Pipelining

# Gotta Do Laundry

- Avi, Bora, Caroline, Dan each have one load of clothes to wash, dry, fold, and put away
  - Washer takes 30 minutes
  - Dryer takes 30 minutes
  - “Folder” takes 30 minutes
  - “Stasher” takes 30 minutes to put clothes into drawers

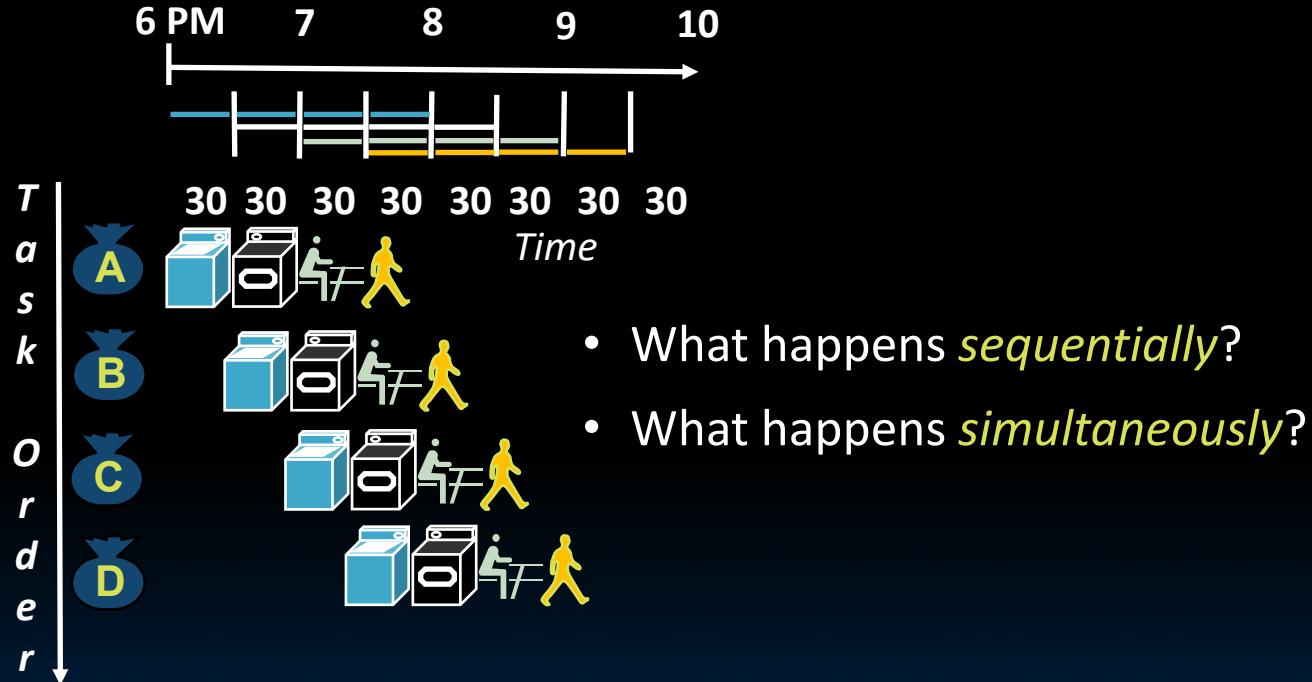


# Sequential Laundry



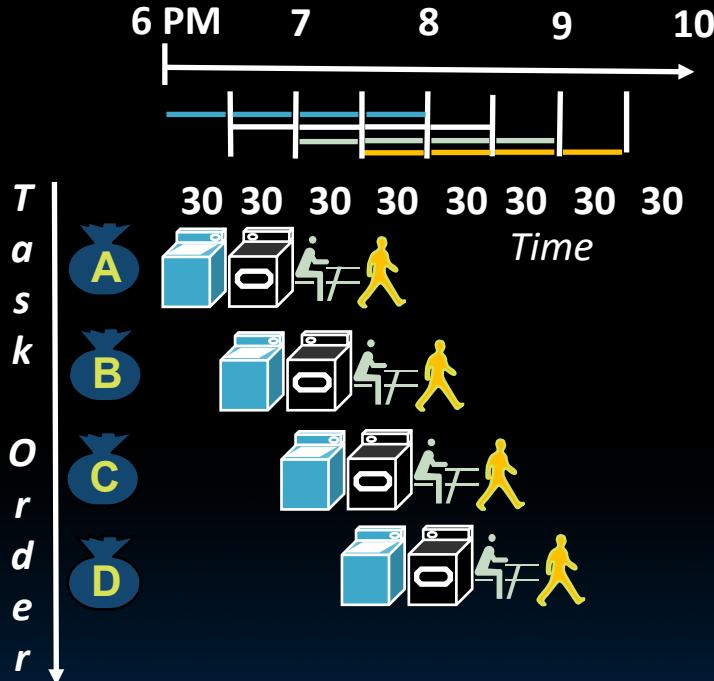
Sequential laundry takes 8 hours for 4 loads!

# Pipelined Laundry



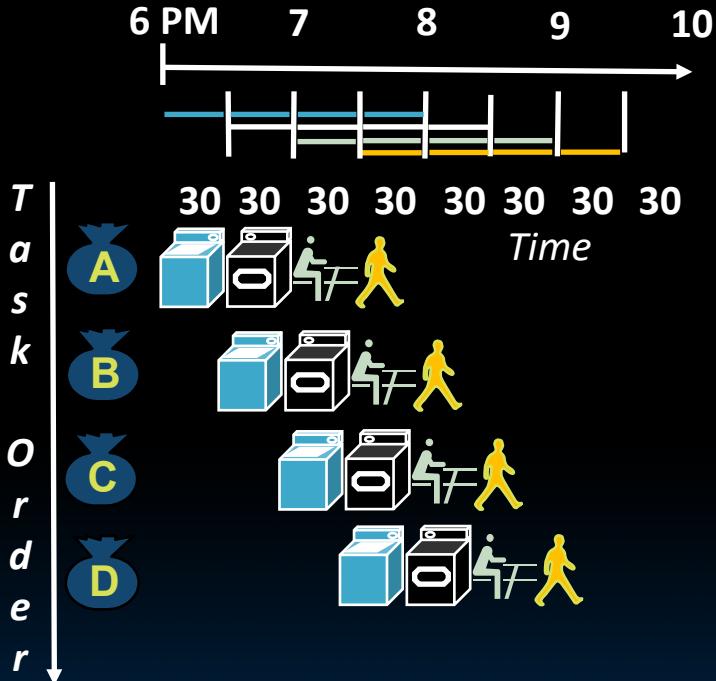
Pipelined laundry takes 3.5 hours for 4 loads!

# Sequential Laundry



- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- **Multiple** tasks operating simultaneously using different resources
- Potential speedup = **Number of pipe stages**
- Time to “fill” pipeline and time to “drain” it reduces speedup:  
2.3X v. 4X in this example

# Sequential Laundry



Suppose:

- new Washer takes 20 minutes
- new Stasher takes 20 minutes.

How much faster is pipeline?

Pipeline rate limited by **slowest** pipeline stage

Unbalanced lengths of pipe stages reduce speedup